

Definition: Private-key Encryption Scheme: Specify a message space \mathcal{M} , and Gen , Dec , and Enc algorithms. Gen is a probabilistic algorithm that outputs a key k . Enc takes k and $m \in \mathcal{M}$, and outputs ciphertext c . Notation: $c = \text{Enc}_k(m)$. Dec takes k and c , and outputs m . Notation: $m = \text{Dec}_k(c)$. Must have $\text{Dec}_k(\text{Enc}_k(m)) = m$ for all k . The set of valid keys is \mathcal{K} . WLOG, assume Gen chooses a uniform $k \in \mathcal{K}$.

Definition: Kerchoffs' Principle: An encryption scheme should be designed to be secure *even if* an eavesdropper knows all the details of the scheme, so long as the attacker doesn't know the key being used.

Definition: Sufficient Key-space Principle: Any secure encryption scheme must have a key space that is sufficiently large to make an exhaustive-search attack infeasible. Necessary, but not sufficient.

Theorem: Bayes Theorem: $\Pr[A|B] = \frac{\Pr[B|A] \cdot \Pr[A]}{\Pr[B]}$.

Definition: Perfect Secrecy: An encryption scheme $(\text{Gen}, \text{Enc}, \text{Dec})$ with message space \mathcal{M} such that for every probability distribution over \mathcal{M} , every message $m \in \mathcal{M}$, and every ciphertext $c \in \mathcal{C}$, for which $\Pr[C = c] > 0$, $\Pr[M = m | C = c] = \Pr[M = m]$. Equivalently, $\forall m, m' \in \mathcal{M}, c \in \mathcal{C}$, $\Pr[\text{Enc}_K(m) = c] = \Pr[\text{Enc}_K(m') = c]$.

Definition: $\text{PrivK}_{\mathcal{A}, \Pi}^{\text{eav}}$: (Adversarial Indistinguishability Experiment): The Adversary \mathcal{A} outputs $m_0, m_1 \in \mathcal{M}$. $k \leftarrow \text{Gen}$, $b \in \{0, 1\}$ uniformly. $c \leftarrow \text{Enc}_k(m_b)$ is given to \mathcal{A} , called challenge ciphertext. $b' \leftarrow \mathcal{A}$. Output is 1 ("success") iff $b' = b$, notated $\text{PrivK}_{\mathcal{A}, \Pi}^{\text{eav}} = 1$.

Definition: Perfect Indistinguishability: An encryption scheme $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ with \mathcal{M} , such that $\forall \mathcal{A}$, $\Pr[\text{PrivK}_{\mathcal{A}, \Pi}^{\text{eav}} = 1] = \frac{1}{2}$.

Theorem: Perfect Indistinguishability \Leftrightarrow Perfect Secrecy.

Definition: One-time Pad: Fix $\ell > 0$. $\mathcal{M} = \mathcal{K} = \mathcal{C} = \{0, 1\}^\ell$ (binary strings length ℓ). Gen : uniform $k \in \mathcal{K}$. Enc : $c = k \oplus m$, \oplus is bitwise xor. Dec : $m = k \oplus c$.

Theorem: The one-time pad encryption scheme is perfectly secret.

Theorem: In a perfectly secure encryption scheme, $|\mathcal{K}| \geq |\mathcal{M}|$. ($|X|$ denotes magnitude/size of X .)

Theorem: Shannon's Theorem: Let $(\text{Gen}, \text{Enc}, \text{Dec})$ be an encryption scheme with $|\mathcal{M}| = |\mathcal{K}| = |\mathcal{C}|$. It is perfectly secret iff all $k \in \mathcal{K}$ are chosen with probability $1/|\mathcal{K}|$ by Gen , and $\forall m \in \mathcal{M}, c \in \mathcal{C}$, $\exists k \in \mathcal{K}$ such that $c = \text{Enc}_k(m)$.

Definition: A cryptographic scheme is (t, ϵ) -secure if any adversary running for time at most t succeeds in breaking the scheme with probability at most ϵ .

Definition: PPT (Probabilistic Polynomial Time): An adversary which runs for time at most $p(n)$, where n is the security parameter (length of key), and p is a polynomial.

Definition: negl (Negligible): A function f from the natural numbers to the non-negative real numbers such that for every positive polynomial p there is an $N \in \mathbb{N}$ such that $\forall n > N$, $f(n) < \frac{1}{p(n)}$.

Definition: Secure: A scheme where any PPT adversary succeeds in breaking the scheme with at most negligible probability.

Definition: Probabilistic: An algorithm that can "toss a coin" - access unbiased random bits - as necessary.

Theorem: Let $\text{negl}_1, \text{negl}_2$ be negligible functions, p a polynomial. Then $\text{negl}_1(n) + \text{negl}_2(n)$ and $p(n) \cdot \text{negl}_1(n)$ are both negligible.

Definition: Secure: A scheme for which every PPT adversary \mathcal{A} carrying out an attack of some formally specified type, the probability that \mathcal{A} succeeds is negligible.

Definition: We denote an error from Dec by \perp (bottom), when it is asked to decrypt a non-valid ciphertext.

Definition: Fixed-length encryption scheme: An encryption scheme such that for a $k \leftarrow \text{Gen}(1^n)$, Enc_k is only defined for messages $m \in \{0, 1\}^{\ell(n)}$ (fixed length messages).

Definition: $\text{PrivK}_{\mathcal{A}, \Pi}^{\text{eav}}$ (n) (Adversarial Indistinguishability Experiment-EAV): , where n is the security parameter, and success is defined as before. However, \mathcal{A} is a PPT adversary, $|m_0| = |m_1|$, but the guessing for $b = b'$ is identical.

Definition: EAV-secure (indistinguishable encryptions in the presence of an eavesdropper): A private key encryption scheme $(\text{Gen}, \text{Enc}, \text{Dec})$ such that for all PPT adversaries \mathcal{A} , for all n , $\Pr[\text{PrivK}_{\mathcal{A}, \Pi}^{\text{eav}}(n) = 1] \leq \frac{1}{2} + \text{negl}(n)$, where the probability is taken over the randomness used by \mathcal{A} and the encryption scheme Π . Equivalently: $|\Pr[\text{out}_{\mathcal{A}}(\text{PrivK}_{\mathcal{A}, \Pi}^{\text{eav}}(n, 0)) = 1] - \Pr[\text{out}_{\mathcal{A}}(\text{PrivK}_{\mathcal{A}, \Pi}^{\text{eav}}(n, 1))]| \leq \text{negl}(n)$.

Theorem: Let $\Pi = (\text{Enc}, \text{Dec})$ be a fixed-length private-key encryption scheme for messages of length ℓ that has indistinguishable encryptions in the presence of an eavesdropper. Then for all PPT adversaries \mathcal{A} and any $i \in \{1, \dots, \ell\}$, there is a negligible function negl such that $\Pr[\mathcal{A}(1^n, \text{Enc}_k(m)) = m^i] \leq \frac{1}{2} + \text{negl}(n)$, where m^i is the i^{th} bit of m .

Theorem: Let (Enc, Dec) be a fixed-length private key encryption scheme for messages of length ℓ that is EAV-secure. Then for any PPT algorithm \mathcal{A} there is a PPT algorithm \mathcal{A}' such that for any $S \subseteq \{0, 1\}^\ell$ and any function $f: \{0, 1\}^\ell \rightarrow \{0, 1\}$, $|\Pr[\mathcal{A}(1^n, \text{Enc}_k(m)) = f(m)] - \Pr[\mathcal{A}(1^n) = f(m)]| \leq \text{negl}(n)$. That is, \mathcal{A} cannot determine any function f of the original message m , given the ciphertext, with more than negligible probability better than when not given the ciphertext.

Definition: PRG (Pseudo-Random Generator): Let ℓ be a polynomial and G be a deterministic polynomial-time algorithm such that for any n and any input $s \in \{0, 1\}^n$, the result $G(s)$ is a string of length $\ell(n)$. The following must hold: For every n , $\ell(n) > n$. For any PPT algorithm D , there is a negligible function negl such that $|\Pr[D(G(s)) = 1] - \Pr[D(r) = 1]| \leq \text{negl}(n)$. ℓ is the expansion factor of G .

Construction: Stream Cipher: Let G be a pseudorandom generator with expansion factor ℓ . Let $\text{Gen}(1^n)$ output a uniform $k \in \{0, 1\}^n$. Let $c = \text{Enc}_k(m) = G(k) \oplus m$. Let $\text{Dec}_k(c) = G(k) \oplus c$. This is an EAV-secure private-key encryption scheme.

Definition: $\text{PrivK}_{\mathcal{A}, \Pi}^{\text{mult}}$: The EAV experiment, except \mathcal{A} presents 2 equal length lists of equal length messages, $\vec{M}_0 = (m_{0,1}, \dots, m_{0,\ell})$ and $\vec{M}_1 = (m_{1,1}, \dots, m_{1,\ell})$, the challenger chooses one of the lists and returns the ciphertext of all messages from that list, and \mathcal{A} attempts to determine which list was chosen.

Definition: Multiple-EAV-Secure: Same as EAV secure, except with the $\text{PrivK}_{\mathcal{A}, \Pi}^{\text{mult}}$ experiment.

Theorem: There are private-key encryption schemes which are EAV-secure but not multiple-EAV-secure.

Theorem: Any multiple-EAV-secure private-key encryption scheme is also EAV-secure.

Theorem: If Π is a stateless encryption scheme in which Enc is deterministic, then Π cannot be multiple-EAV-secure.

Definition: $\text{PrivK}_{\mathcal{A}, \Pi}^{\text{CPA}}(n)$: $k \leftarrow \text{Gen}(1^n)$, then adversary \mathcal{A} is given 1^n and oracle access to $\text{Enc}_k(\cdot)$. \mathcal{A} , after making its oracle calls, outputs m_0, m_1 , a pair of same length messages. b is chosen, $c = \text{Enc}_k(m_b)$ is computed and returned, and \mathcal{A} outputs b' . \mathcal{A} "succeeds" if $b' = b$, and the experiment outputs 1. Else, the experiment outputs 0.

Definition: CPA-secure (Chosen Plaintext Attack): A private-key encryption scheme $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ such that for all PPT adversaries \mathcal{A} , $\Pr[\text{PrivK}_{\mathcal{A}, \Pi}^{\text{CPA}}(n) = 1] \leq \frac{1}{2} + \text{negl}(n)$.

Definition: $\text{PrivK}_{\mathcal{A}, \Pi}^{\text{LR-CPA}}(n)$: Same as $\text{PrivK}_{\mathcal{A}, \Pi}^{\text{mult}}$, but for CPA-security. Extends to Multiple-CPA-security.

Theorem: CPA-secure \Rightarrow multiple-CPA-secure.

Definition: Keyed Function: A function $F: \{0, 1\}^* \times \{0, 1\}^* \rightarrow \{0, 1\}^*$, with first input called key k . It is efficient if there is a polynomial-time algorithm that computes $F(k, x)$ given k and x .

Definition: Length-Preserving: A keyed function such that $\ell_{\text{key}}(n) = \ell_{\text{in}}(n) = \ell_{\text{out}}(n)$.

Definition: Func_n: The set of all functions mapping n -bit strings to n -bit strings. $|\text{Func}_n| = 2^{n \cdot 2^n}$.

Definition: PRF (Pseudo-Random Function): An efficient, length-preserving keyed function f such that for all PPT distinguishers D , $|\Pr[D^{F_k(\cdot)}(1^n) = 1] - \Pr[D^{f(\cdot)}(1^n) = 1]| \leq \text{negl}(n)$, where f is chosen uniformly from Func_n . $|\text{PRF}| = 2^n$, there are at most that many distinct functions. (Some are not secure.)

Definition: Permutation: A keyed function F such that $\ell_{\text{in}} = \ell_{\text{out}}$, and for all $k \in \mathcal{K}$, $F_k: \{0, 1\}^{\ell_{\text{in}}(n)} \rightarrow \{0, 1\}^{\ell_{\text{out}}(n)}$ is one-to-one. F_k is efficient if $F_k(x)$ and $F_k^{-1}(x)$ are computable with a polynomial-time algorithm.

Definition: PRP (Pseudo-Random Permutation): Same as a PRF, except F must be indistinguishable from a random $f \in \text{Perm}_n$, the set of truly random permutations.

Theorem: If F is a PRP and $\ell_{\text{in}}(n) \geq n$, F is a PRF.

Definition: Strong PRP: Let $F: \{0, 1\}^* \times \{0, 1\}^* \rightarrow \{0, 1\}^*$ be an efficient, length-preserving, keyed permutation such that for all PPT distinguishers D , $|\Pr[D^{F_k(\cdot), F_k^{-1}(\cdot)}(1^n) = 1] - \Pr[D^{f(\cdot), f^{-1}(\cdot)}(1^n) = 1]| \leq \text{negl}(n)$, where $f \in \text{Perm}_n$ uniformly. Any strong PRP is a PRP.

Definition: Synchronized: Stream cipher mode where sender and receiver must know how much plaintext has been encrypted/decrypted so far. Typically used in a single session between parties.

Definition: Unsynchronized: Stream cipher mode which is stateless, taking a new IV each time.

Definition: ECB Mode (Electronic Code Book): Here, $c := \langle F_k(m_1), F_k(m_2), \dots, F_k(m_\ell) \rangle$, where $m = m_1, m_2, \dots, m_\ell$ is the message and F is a block cipher of length n . Deterministic, and therefore not CPA-secure. Should not be used, only included for historical significance.

Definition: CBC Mode (Cipher Block Chaining): Choose an IV of length n . Then, $c_0 = IV$, $c_1 = F_k(c_0 \oplus m_1)$, $c_2 = F_k(c_1 \oplus m_2)$, and so on. To decrypt, compute $m_\ell = F_k^{-1}(c_\ell) \oplus c_{\ell-1}$, $m_{\ell-1} = F_k^{-1}(c_{\ell-1}) \oplus c_{\ell-2}$, and so on. If F is a PRP, and IV is chosen uniformly at random, then CBC mode is CPA-secure. It cannot be computed in parallel, since encrypting c_i requires c_{i-1} for $i > 0$. Using c_ℓ as IV for the next encryption is not secure.

Definition: OFB Mode (Output FeedBack): Let IV be uniformly chosen of length n . Then $c_0 = IV$, $c_1 = F_k(IV) \oplus m_1$, $c_2 = F_k(F_k(IV)) \oplus m_2$, \dots , $c_\ell = F_k^\ell(IV) \oplus m_\ell$, where F_k^ℓ denotes F_k applied ℓ times. F need not be invertible, and m_ℓ need not be of length n , the message may be truncated to match its length. OFB mode is CPA-secure. Using $F_k^\ell(IV)$ as the next IV , producing a synchronized stream cipher, it remains secure.

Definition: CTR Mode (Counter): Pick an $\text{ctr} = IV$, then $c_0 = \text{ctr}$, $c_1 = F_k(\text{ctr} + 1) \oplus m_1$, \dots , $c_\ell = F_k(\text{ctr} + \ell) \oplus m_\ell$. F need not be invertible. Here, the encryption can be fully parallelized. CTR mode is CPA-secure, assuming F is a PRF.

The stateful variant, where $F_k(\text{ctr} + \ell)$ is used as the new IV , remains secure.

Note: None of these schemes achieve message integrity in the sense of chapter 4.

Definition: $\text{PrivK}_{\mathcal{A}, \Pi}^{\text{CCA}}(n)$: The adversary \mathcal{A} is given access to a decryption oracle in addition to an encryption oracle, then outputs m_0, m_1 , gets $c := \text{Enc}_k(m_b)$, the challenge ciphertext, and tries to determine b' . \mathcal{A} again has oracle access, but cannot query the decryption oracle with c . Success is as defined in previous experiments.

Definition: CCA-Secure (Chosen Ciphertext Attack): A private-key encryption scheme Π such that for all PPT adversaries \mathcal{A} , $\Pr[\text{PrivK}_{\mathcal{A}, \Pi}^{\text{CCA}}(n) = 1] \leq \frac{1}{2} + \text{negl}(n)$. Any CCA-secure scheme is also multiple-CCA-secure.

Note: NOTHING above this point is CCA-secure.

Definition: Non-Malleability: An encryption scheme with the property that if the adversary tries to modify a given ciphertext, the result is either an invalid ciphertext or one whose corresponding plaintext has no relation to the original plaintext.

Definition: MAC (Message Authentication Code): Three probabilistic polynomial-time algorithms GenMacVrfy such that Gen takes 1^n , outputs k with $|k| \geq n$. Mac , the tag-generation algorithm, takes k and $m \in \{0, 1\}^*$ and outputs tag t . Deterministic Vrfy takes k, m, t , and outputs b , where $b = 1$ means t is a valid tag for m with key k , and $b = 0$ means it is not. It must be that $\text{Vrfy}_k(m, \text{Mac}_k(m)) = 1$. If Mac_k is only defined for $m \in \{0, 1\}^{\ell(n)}$, we call it a fixed-length MAC.

Definition: Canonical Verification: Deterministic MACs (Mac is deterministic), where $\text{Vrfy}_k(m, t)$ computes $\tilde{t} := \text{Mac}_k(m)$, and outputs 1 iff $\tilde{t} = t$.

Definition: $\text{Mac-forge}_{\mathcal{A}, \Pi}(n)$: $k \leftarrow \text{Gen}(1^n)$. \mathcal{A} is

given 1^n and oracle access to $\text{Mac}_k(\cdot)$. Eventually outputs (m, t) . Success is defined as $\text{Vrfy}(m, t) = 1$, and m is not a message previously queried from the oracle.

Definition: Secure MAC (Existentially Unforgeable Under an Adaptive Chosen-Message Attack): A MAC such that for all PPT adversaries \mathcal{A} , $\Pr[\text{Mac-forge}_{\mathcal{A}, \Pi}(n) = 1] \leq \text{negl}(n)$. Note: This definition offers no protection against replay attacks.

Definition: Strong MAC: MAC such that (m, t) cannot have been previously output by the oracle, but using (m, t') is a valid guess.

Theorem: With a canonical Vrfy , Strong Mac \Leftrightarrow Secure MAC.

Construction: : $\text{Mac}_k(m) = F_k(m)$, where $k \in \{0, 1\}^n$, $m \in \{0, 1\}^n$, and F a PRF. Vrfy is canonical. If $|m| \neq |k|$, Mac outputs nothing, and Vrfy outputs 0. This construction is a secure fixed-length MAC.

Construction: Another Mac_k : Chop m into n -bit blocks, m_1, m_2, \dots, m_d , let $t_i = \text{Mac}'(m_i)$, and use (t_1, t_2, \dots, t_d) as the tag.

This is bad. This can be easily broken using a reordering attack. Present $m = m_1, m_2$, get tag t_1, t_2 . Then message $m' = m_2, m_1$ will have tag t_2, t_1 , which will pass Vrfy .

To combat this attack, break m into $\frac{n}{2}$ -bit blocks m_1, \dots, m_d , then $t_i = \text{Mac}'_k(\langle i \rangle \parallel m_i)$, where $\langle i \rangle$ is the $\frac{n}{2}$ -bit binary encoding of i . This prevents the reordering attack.

This scheme is still insecure. Since we have an arbitrary-length message Mac , we can use a truncation attack, and present $m = m_1, m_2, m_3$, get (t_1, t_2, t_3) . Then we can present $m' = m_1, m_2$. The tag (t_1, t_2) will be valid for m' .

To prevent the truncation attack, we will include the length ℓ of the full message in the calculation. We will chop our message into $\frac{n}{3}$ -bit blocks. Then $t_i = \text{Mac}'_k(\langle \ell \rangle \parallel \langle i \rangle \parallel m_i)$. Note: We pad the last block with 0's if necessary. The tag will be (t_1, \dots) . By this point, we are sending 4ℓ bits.

Unfortunately, even this scheme is still insecure. It can be attacked with a “mix and match” attack. For example, get tag $t = (t_1, t_2, t_3)$ for $m = m_1, m_2, m_3$. Take another message, same length, $m' = m_4, m_5, m_6$, get tag $t' = (t_4, t_5, t_6)$. Then (t_1, t_5, t_6) is a valid tag for m_1, m_5, m_6 , which has never been queried from the oracle before.

Finally, let's fix all of this! We'll chop our message $m = m_1, \dots, m_d$ into $\frac{n}{4}$ bit blocks, and pick a random $\frac{n}{4}$ -bit value r for the entire message, and $t_i = \text{Mac}'_k(r \parallel \langle \ell \rangle \parallel \langle i \rangle \parallel m_i)$. $\text{Mac}_k(m) = (r, t_1, \dots, t_d)$. At this point, this is not a deterministic Mac , so Vrfy has to behave slightly differently, taking into account the random r passed to it. It can reconstruct the tag as above, with this slight extra step.

This is secure!

But it does produce a tag of 4x the length of the message.

Construction: CBC-MAC: Used widely in practice. On input a key $k \in \{0, 1\}^n$, m of length $\ell(n) \cdot n$, let $\ell = \ell(n)$, parse $m = m_1, \dots, m_\ell$, set $t_0 := 0^n$, then for $i \in \{1, \ell\}$, $t_i := F_k(t_{i-1})$, where F is a PRF. Output t_ℓ only as the tag. Vrfy is done in the canonical way.

To extend this to arbitrary length messages, *prepend* the message with length $|m|$, encoded as an n -bit string.

Alternatively, have keys k_1, k_2 , compute CBC-MAC using k_1 , then output tag $\hat{t} := F_{k_2}(t)$.